


V1 - a single agent, decides for itself when and what to search

1 agent - the graph is a single node: START -> agent -> END



Asymmetric-trust rule: HIGH evidence wins by default, overturned only when reasoning names a specific finding the evidence missed; MEDIUM/LOW is set aside entirely, decided purely from the model's own reasoning.

purple = LLM call - teal = retrieval (no LLM, the model decides for itself whether to call it)