


V2 - understand -> (search || reasoning) -> answer

4 agents, 4 LLM calls/item - no verifier



Node 2 and Node 3 never see each other's output; Node 4 applies an explicit rule, not free-form synthesis.

Different from the multi-expert panel design (MedAgents) tried earlier - this is 4 steps of ONE pipeline, not several opinions debating.

purple = LLM call - teal = retrieval (no LLM, called directly by the graph - deterministic)