


V3 - understand -> (search || reasoning) -> answer + long-term memory

4 agents, 4 LLM calls/item - no verifier; decider recalls/remembers lessons



decider recalls lessons from similar earlier cases before answering, then writes one back - no separate agent or LLM call, just two store reads/writes around the existing decider call.

purple = LLM call - teal = retrieval (no LLM) - amber = long-term memory (SQLite, no LLM)